

420+ DevOps Interview Questions



Learn With Shubham Praharaj



DevOps Interview Master Guide

(Structured Questions Bank + What Interviewers Expect)

This document is a **complete DevOps interview question bank** designed to help you prepare for roles like:

- DevOps Engineer
- Cloud DevOps Engineer
- Platform Engineer
- Site Reliability Engineer (SRE)
- Kubernetes / Terraform / CI-CD Engineer

It is structured by **skills + difficulty level + production scenarios**, so you can revise quickly and answer confidently in interviews.

How to Use This Guide (Fast Preparation Plan)

Step 1 — First Pass (Quick Scan)

Try answering each question in **60–90 seconds**.

Step 2 — Second Pass (Weak Areas Only)

Mark the ones where you struggle and revise again.

Step 3 — Scenario Practice

Focus on scenario-based questions (these are asked in senior roles).

- Best interview structure to answer **any question**:

Concept → Why → How it works → Example → Best Practices → Troubleshooting

Skills Covered

- AWS DevOps
 - SRE
 - Kubernetes
 - Jenkins
 - Terraform
 - Ansible
 - Shell & Python Automation
 - Monitoring, Logging & Incident Response
 - Security & DevSecOps
-



Table of Contents

Section 1 — AWS DevOps Interview Questions (50)

- Basic (1–10)
- Advanced (11–20)
- Intermediate (21–30)
- Expert (31–40)
- Production Scenarios (41–50)

Section 2 — SRE Interview Questions (60+)

- Beginner
- Intermediate
- Situation-Based / Incident Questions

Section 3 — Kubernetes Interview Questions (60)

- Beginner

- Intermediate
- Experienced

Section 4 — Jenkins Interview Questions (65)

- Beginner
- Intermediate / Advanced
- Expert

Section 5 — Terraform Interview Questions (70+)

- Beginner
- Intermediate
- Expert

Section 6 — Ansible Interview Questions (40)

- Basics
- Inventory
- Playbooks & Troubleshooting
- Variables & Modules
- Advanced

Section 7 — Shell & Python Interview Questions (50)

- Beginner to Expert
- Production Scenarios

Section 1 — AWS DevOps Interview Questions (50)

Basic Conceptual Level (1–10)

1. Explain the core principles of DevOps and its benefits in AWS

Answer (What interviewer expects):

DevOps improves delivery speed and stability by promoting:

- Automation (CI/CD + IaC)
- Continuous feedback
- Collaboration between Dev + Ops
- Monitoring-driven improvement

In AWS, DevOps becomes faster because AWS provides managed services like: CloudWatch, CodePipeline, Auto Scaling, ECS/EKS, Lambda.

2. IaC vs IaaS

Answer:

- **IaaS:** AWS provides infrastructure resources like EC2, EBS, VPC.
 - **IaC:** You manage infra using code (Terraform/CloudFormation), enabling reproducibility.
-

3. 3 main service categories in AWS

Answer:

- **Compute** (EC2, Lambda, ECS/EKS)
 - **Storage** (S3, EBS, EFS)
 - **Networking** (VPC, Route53, CloudFront)
-

4. Types of EC2 instances and choosing the right one

Answer:

EC2 instance families include:

- General purpose (balanced)
- Compute optimized (CPU heavy)
- Memory optimized (in-memory workloads)
- Storage optimized (high IOPS)
- GPU instances (ML/graphics)

Choice depends on workload: CPU usage, memory needs, traffic pattern, and budget.

5. Security Groups vs NACLs

Answer:

- **Security Group:** stateful firewall at instance level
- **NACL:** stateless firewall at subnet level

Security groups allow rules only, NACL supports allow + deny.

6. Benefits of VPC

Answer:

VPC provides:

- network isolation
 - control of subnets and routing
 - security using SG/NACL
 - private connectivity (VPC endpoints)
 - hybrid integration (VPN/Direct Connect)
-

7. S3 storage classes

Answer:

Common S3 classes:

- Standard
- Standard-IA
- One Zone-IA
- Glacier Instant Retrieval
- Glacier Flexible Retrieval
- Deep Archive

Used based on access frequency and retention.

8. CloudWatch usage

Answer:

CloudWatch provides:

- Metrics

- Logs
- Alarms
- Dashboards
- Events (EventBridge)

Used for monitoring CPU, latency, errors, and sending alerts.

9. AWS Lambda use-case

Answer:

Lambda runs code without managing servers.

Best for:

- event-based systems
 - automation
 - APIs
 - scheduled tasks
 - lightweight microservices
-

10. Autoscaling in AWS

Answer:

Auto Scaling increases or decreases compute capacity based on metrics like: CPU, memory, request count, queue length.

Implemented using:

- Auto Scaling Groups + Launch templates
- Target tracking policies

Advanced Conceptual Level (11–20)

11. CodePipeline vs CodeDeploy

Answer:

- **CodePipeline:** end-to-end CI/CD orchestration
- **CodeDeploy:** deployment engine (EC2, Lambda, ECS)

CodePipeline manages stages, CodeDeploy handles deployment strategies.

12. Serverless architecture: benefits + challenges

Answer:

Benefits:

- no servers to manage
- scales automatically
- cost-efficient for event workloads

Challenges:

- cold start latency
- debugging complexity
- vendor lock-in
- tracing distributed flows is harder

13. ECS vs EKS

Answer:

- ECS = simpler AWS-managed container orchestration
- EKS = Kubernetes-based (more flexibility, industry standard)

Choose ECS for simplicity, EKS for K8s ecosystem + portability.

14. Role of Terraform/CloudFormation in DevOps

Answer:

They implement Infrastructure as Code enabling:

- repeatable deployments
 - reduced manual errors
 - version-controlled infra changes
 - faster environment creation
-

15. IaC testing approach

Answer:

IaC testing validates:

- syntax and formatting (linting)
- security checks (policy as code)
- unit/integration validation (test environments)

Tools: terraform validate, tflint, checkov, terratest.

16. Disaster recovery strategies in AWS

Answer:

- Backup & restore
- Pilot light
- Warm standby
- Active-active multi-region

Chosen based on RTO and RPO requirements.

17. IAM + VPC security best practices

Answer:

- least privilege roles
 - avoid access keys (use IAM roles)
 - enable MFA
 - use private subnets + NAT
 - use VPC endpoints
 - logging with CloudTrail + Config
-

18. AWS cost optimization strategies

Answer:

- right-size EC2
- Spot and Savings Plans
- S3 lifecycle policies
- turn off unused resources

- set budgets + alerts
 - use load testing to avoid overprovisioning
-

19. Observability for serverless

Answer:

Use:

- CloudWatch Logs Insights
 - X-Ray tracing
 - OpenSearch
 - structured logging + correlation IDs
-

20. Blue/Green deployment in AWS

Answer:

Blue/Green deployments reduce downtime by running:

- Blue = current stable environment
- Green = new version

Traffic switches after validation using ALB or CodeDeploy.



Intermediate Level (21–30)

21. Explain a DevOps project you worked on

Answer format:

- problem statement
 - tools used
 - pipeline flow
 - deployment strategy
 - monitoring + incident results
 - what you improved
-

22. Infrastructure changes with minimal downtime

Answer:

Use:

- rolling updates
 - blue/green
 - canary releases
 - feature flags
 - database migrations carefully (backward compatible schema)
-

23. Ansible/Chef automation experience

Answer:

Use configuration management for:

- OS patching
- installing packages
- configuring services

- environment consistency
 - repeatability across dev/stage/prod
-

24. Troubleshooting AWS deployments

Answer approach:

Check in this order:

- CloudWatch alarms/logs
 - deployment pipeline logs
 - IAM permissions
 - security group/network routing
 - resource limits/throttling
-

25. Monitoring AWS applications

Answer:

Use metrics:

- latency, error rate, throughput
- CPU/memory
- database connections
- queue depth

Tools: CloudWatch, X-Ray, OpenTelemetry, dashboards.

26. Experience with IaC scripts

Answer:

Mention:

- modules
 - remote backend/state locking
 - code reviews
 - environment separation (dev/stage/prod)
 - drift detection
-

27. Kubernetes usage in AWS

Answer:

EKS used for microservices:

- deployments
 - ingress
 - autoscaling
 - secrets/configmaps
 - logging/monitoring integration
-

28. CI/CD pipelines in AWS

Answer:

Typical flow:

Git → build → test → package → deploy → verify → rollback plan

Tools: CodePipeline, Jenkins, GitHub Actions, ArgoCD.

29. Collaboration with Dev & Security teams

Answer:

- shared pipeline ownership
 - DevSecOps scanning in CI
 - shared runbooks
 - incident review meetings
 - clear approval workflows
-

30. Incident response experience

Answer:

- detect → triage → mitigate → fix → postmortem
 - rollback strategy
 - root cause analysis
 - preventive automation
-
-



Expert Level (31–40)

31. Advanced AWS services usage

Examples:

- Step Functions for workflows
 - Lambda Layers for shared dependencies
 - custom CloudFormation resources
-

32. Encryption for sensitive data

Answer:

Use:

- KMS
 - Secrets Manager
 - TLS everywhere
 - encryption at rest for EBS/RDS/S3
 - key rotation policies
-

33. Serverless security best practices

Answer:

- least privilege IAM role
 - restrict Lambda network with VPC
 - validate inputs
 - API Gateway rate limiting
 - WAF integration
 - logging + anomaly detection
-

34. HA and scalable web architecture

Answer:

Multi-AZ design:

- ALB

- Auto Scaling
 - RDS Multi-AZ or Aurora
 - CloudFront + caching
 - disaster recovery plan
-

35. AWS performance optimization

Answer:

- caching (CloudFront, ElastiCache)
 - reduce DB calls
 - asynchronous design (SQS/Kafka)
 - tune instance types and scaling triggers
-

36. Automated compliance & audits

Answer:

Use:

- AWS Config rules
 - Security Hub
 - GuardDuty
 - IaC policy checks (OPA, Checkov)
-

37. Staying updated

Answer:

- AWS What's New
 - architecture blogs and documentation
 - hands-on labs
 - internal knowledge sharing
-

38. Hard technical problem you solved

Best structure:

- impact
 - root cause
 - fix
 - prevention
 - lessons learned
-

39. Cost management strategy

Answer:

- cost explorer analysis
- tagging enforcement
- scheduled shutdown
- rightsizing
- savings plans
- reserved capacity for steady workloads

40. Building DevOps culture

Answer:

- encourage automation
 - reduce blame culture
 - shared responsibility
 - documentation/runbooks
 - standardization across teams
-
-

Expert Production Scenarios (41–50)

41. Flash sale traffic surge causing outages

Answer:

- validate ALB target health
 - check scaling triggers
 - enable caching + CloudFront
 - scale DB read replicas
 - use queue for async traffic
 - rollback quickly if release caused failure
-

42. Production database corrupted

Answer:

- stop writes
 - restore from snapshot
 - use point-in-time recovery (PITR)
 - validate integrity
 - introduce backup policy + deletion protections
-

43. Slow build times in CI/CD

Answer:

- parallelize builds
 - enable caching dependencies
 - use dedicated agents
 - optimize tests (unit vs integration separation)
 - remove unnecessary pipeline steps
-

44. Security vulnerability in public API

Answer:

- immediate mitigation (WAF rule, rate-limit, block attack vector)
- patch + deploy fix
- rotate credentials
- audit logs
- postmortem + prevention

45. Migrate legacy app to AWS

Answer:

- assessment (dependencies, DB, traffic)
 - choose migration strategy (rehost/refactor)
 - build staging infra using IaC
 - testing + cutover plan
 - rollback plan
-

46. High latency issue

Answer:

- check CloudWatch metrics
 - ALB response time
 - DB slow queries
 - external API latency
 - caching improvements + scaling
-

47. Unauthorized S3 access attempt

Answer:

- check CloudTrail logs
- block access via policy

- enforce MFA delete
 - enable versioning
 - rotate credentials and review IAM policies
-

48. Automate microservices deployment

Answer:

CI/CD design:

- build Docker images → push registry
 - deploy to Kubernetes/ECS
 - canary + rollback
 - integrate tests + security scan
 - add monitoring gates
-

49. AWS cost is too high

Answer:

- identify top cost services
 - remove unused resources
 - rightsizing + Spot
 - S3 lifecycle
 - database optimization
 - enforce budgets + alerts
-

50. Adopting DevOps culture

Answer:

- start with CI + automation
- build shared pipelines
- define ownership
- improve collaboration between Dev/Ops
- introduce monitoring + fast feedback

SRE (SITE RELIABILITY ENGINEERING) — INTERVIEW QUESTIONS + ANSWERS (70+ Qs)

Beginner → Intermediate → Production Scenarios | Google Docs Ready Format

How to Use This SRE Section

-  **Round 1:** Answer in 1–2 lines (quick recall)
 -  **Round 2:** Answer in 30–60 seconds (interview mode)
 -  **Round 3:** Practice scenario questions like incident calls
-

PART 1 — SRE ABSOLUTE BEGINNER QUESTIONS (20)

1. What is SRE?

Answer: Site Reliability Engineering (SRE) is a discipline that applies software engineering principles to IT operations to improve reliability, scalability, and efficiency.

2. What is the primary goal of an SRE?

Answer: To ensure systems are **reliable, scalable, and efficient** while balancing the need for **feature velocity vs stability**.

3. What is the difference between DevOps and SRE?

Answer:

- **DevOps:** culture/practices to improve delivery and collaboration
 - **SRE:** implementation model that uses engineering + reliability metrics to run production systems
-

4. What are the key responsibilities of an SRE?

Answer:

- Monitoring and alerting
 - Incident response and root cause analysis
 - Reliability engineering and automation
 - Capacity planning and performance tuning
 - On-call and operational excellence
-

5. What is reliability in SRE terms?

Answer: Reliability means a system behaves as expected under normal and failure conditions, meeting availability and performance targets.

6. What is Availability?

Answer: Availability is the percentage of time a service is operational and accessible to users (example: 99.9%).

7. What is Latency?

Answer: Latency is the time taken to serve a request (example: response time in milliseconds).

8. What is Throughput?

Answer: Throughput is the number of requests processed per unit time (RPS—requests per second).

9. What is Monitoring and why is it needed?

Answer: Monitoring measures service health through metrics/logs/traces to detect issues early and ensure reliability.

10. What is Alerting?

Answer: Alerting notifies engineers when systems breach defined thresholds (CPU high, error spikes, latency increase).

11. What is Incident Management?

Answer: The structured process of detecting, responding to, resolving, and learning from production incidents.

12. What is On-call in SRE?

Answer: Engineers rotate responsibility to respond to production issues when alerts trigger during or outside work hours.

13. What are runbooks?

Answer: Runbooks are documented step-by-step procedures to handle common incidents (restart service, rollback, clear queue).

14. Why is automation important in SRE?

Answer: Automation reduces human error, speeds up recovery, and improves operational efficiency.

15. What does “infrastructure as code” mean in SRE context?

Answer: Managing infra using version-controlled code to improve consistency, repeatability, and reliability.

16. What is “toil” in SRE?

Answer: Toil is repetitive manual work that produces no long-term value and should ideally be automated.

17. What does it mean to be “production focused”?

Answer: It means prioritizing reliability and customer impact over non-critical improvements.

18. Why are strong problem-solving skills required for SRE?

Answer: SREs must quickly diagnose system failures, reduce downtime, and restore services under pressure.

19. What is the role of collaboration for SRE?

Answer: SREs work with dev/security/product teams to improve system design, release safety, and operational readiness.

20. What is the difference between proactive and reactive work in SRE?

Answer:

- **Reactive:** responding to outages/incidents
 - **Proactive:** preventing incidents (capacity planning, chaos testing, automation)
-
-

PART 2 — SRE INTERMEDIATE QUESTIONS (20)

21. What are SLIs, SLOs, and SLAs?

Answer:

- **SLI (Indicator):** a metric (latency, error rate, availability)
 - **SLO (Objective):** target value for SLIs (99.95% uptime)
 - **SLA (Agreement):** contractual commitment with penalties
-

22. Give examples of SLIs.

Answer:

- Request success rate (2xx)
 - p95/p99 latency
 - Availability (uptime)
 - Queue length
 - CPU/memory saturation
-

23. What is an error budget?

Answer: The allowed amount of failure (downtime/errors) within an SLO period. If exhausted, new feature releases slow down until reliability improves.

24. How do error budgets help decision making?

Answer: They create a measurable balance between shipping features and maintaining stability.

25. What is MTTR?

Answer: Mean Time To Recovery—average time taken to restore service after a failure.

26. What is MTBF?

Answer: Mean Time Between Failures—average time between one failure and the next.

27. What is an RCA (Root Cause Analysis)?

Answer: A structured analysis of why an incident occurred and how to prevent it from recurring.

28. What should a good postmortem include?

Answer:

- What happened (timeline)
 - Impact
 - Root cause
 - Contributing factors
 - Fixes and preventive actions
 - Lessons learned
-

29. What is “blameless postmortem”?

Answer: A postmortem culture focused on improving systems rather than blaming individuals.

30. What is capacity planning?

Answer: Forecasting needed resources (CPU/memory/network/DB capacity) based on traffic growth and workload patterns.

31. What is autoscaling and when is it helpful?

Answer: Automatically scaling resources up/down based on metrics like CPU/RPS. Helpful for variable traffic workloads.

32. What is the difference between vertical and horizontal scaling?

Answer:

- **Vertical:** bigger machine (scale-up)
 - **Horizontal:** more machines (scale-out)
-

33. Explain redundancy and failover.

Answer:

- **Redundancy:** extra components to avoid single points of failure
 - **Failover:** switching traffic to healthy standby resources during failures
-

34. What is High Availability (HA)?

Answer: Designing systems with redundancy so they remain operational even when parts fail.

35. What is disaster recovery (DR)?

Answer: A strategy to restore systems and data after major failures like region outage, data corruption, ransomware, etc.

36. What is RTO and RPO?

Answer:

- **RTO (Recovery Time Objective):** max acceptable downtime
 - **RPO (Recovery Point Objective):** max acceptable data loss window
-

37. What are common DR strategies?

Answer:

- Backup & Restore
 - Pilot Light
 - Warm Standby
 - Active-Active / Multi-region
-

38. What is observability?

Answer: Ability to understand internal system behavior using **metrics, logs, and traces** (plus events).

39. Monitoring vs Observability?

Answer:

- **Monitoring:** detects known problems
 - **Observability:** helps debug unknown problems using deep signals
-

40. What is distributed tracing?

Answer: Tracking a request as it travels through multiple microservices (example tools: Jaeger, AWS X-Ray).

PART 3 — SRE SITUATION-BASED TECHNICAL QUESTIONS (31)

41. A production outage occurs. Walk through your response.

Answer:

1. Acknowledge incident + alert stakeholders
 2. Assess severity and impact
 3. Stop the bleeding (rollback, scale, failover)
 4. Identify root cause (logs/metrics/traces)
 5. Apply permanent fix
 6. Postmortem + prevention actions
-

42. How would you design a CI/CD pipeline for microservices?

Answer:

- Per-service pipelines + versioned artifacts
 - Automated tests (unit → integration → smoke)
 - Security scans (SAST, dependency, container)
 - Deploy to staging, then prod
 - Canary/blue-green rollout
 - Observability + automated rollback
-

43. How would you handle a cloud security breach?

Answer:

1. Contain (block access, isolate systems)
 2. Identify entry vector (IAM, keys, public endpoints)
 3. Rotate secrets + revoke credentials
 4. Collect logs for forensics
 5. Patch vulnerability
 6. Improve guardrails (least privilege, WAF, SIEM alerts)
-

44. Application response time increases significantly. Steps?

Answer:

- Check latency metrics (p95/p99)
 - Identify layer: DNS/CDN → LB → app → DB
 - Look for CPU/memory saturation
 - Check DB query latency / locks
 - Analyze traces/logs
 - Scale or optimize bottleneck
-

45. How do you ensure database high availability in cloud?

Answer:

- Multi-AZ replication
 - Automated backups
 - Read replicas
 - Failover testing
 - Connection pooling
 - Monitor slow queries + locks
-

46. How do you implement blue-green deployment in Kubernetes?

Answer:

- Deploy new version alongside old
 - Route traffic via service selector or ingress weights
 - Validate readiness + metrics
 - Switch traffic fully
 - Rollback fast by routing back
-

47. How do you automate infra provisioning with Terraform?

Answer:

- Use modular Terraform design
 - Remote state + locking
 - CI pipeline runs validate/plan/apply
 - Policy checks (security scanning)
 - Separate workspaces/environments
-

48. You must improve scalability of an application. What do you do?

Answer:

- Identify bottleneck with profiling and metrics
- Add caching (Redis/CDN)
- Scale horizontally

- Optimize DB (indexes, connection pooling)
 - Async workloads using queues
 - Load testing + capacity planning
-

49. How would you integrate security scanning into CI/CD?

Answer:

- SAST scans on pull requests
 - Dependency scanning
 - Container image scanning
 - IaC scanning (Terraform/Helm)
 - Secrets scanning
 - Block merges on critical issues
-

50. How do you manage secrets in production?

Answer:

Use Vault / AWS Secrets Manager / Parameter Store + rotate keys regularly and avoid putting secrets in logs or Git.

51. Deploying a feature requiring changes in multiple services—how do you coordinate?

Answer:

- Backward-compatible APIs
 - Feature flags
 - Versioned contract testing
 - Gradual rollout (canary)
 - Rollback plan ready
 - Communication + deployment checklist
-

52. Describe your disaster recovery plan approach.

Answer:

- Define RTO/RPO
 - Identify critical systems
 - Backup strategy + restore drills
 - Multi-region if required
 - Automate failover steps
 - Run DR tests quarterly
-

53. How do you monitor containerized workloads?

Answer:

- Metrics: CPU/memory, restarts, latency
- Logs: centralized logging
- Traces: distributed tracing

- Alerts: error rate, saturation
 - Dashboard: service health overview
-

54. Migrating a legacy application to cloud—main challenges?

Answer:

- Dependencies and hidden configs
 - Data migration risks
 - Security model shift
 - Networking complexity
 - Performance tuning
 - Downtime avoidance strategy
-

55. How do you optimize cloud costs without impacting reliability?

Answer:

- Right-sizing
- Autoscaling
- Reserved/spot strategy
- Storage lifecycle policies
- Remove unused resources

- Use managed services wisely
 - Monitor cost anomalies
-

56. High CPU alerts in production—how do you handle?

Answer:

- Confirm impact (latency/error increase?)
 - Identify pod(instance consuming CPU)
 - Check traffic patterns
 - Scale horizontally
 - Profile app for loops/memory leak
 - Add caching or optimize queries
-

57. How do you implement automated testing in CI/CD?

Answer:

Unit → Integration → Security → Performance → Smoke tests before production promotion.

58. Network outage affecting cloud connectivity—your response?

Answer:

- Confirm scope (local vs region-wide)
- Failover to alternate region/VPN

- Check routing/DNS/peering
 - Engage provider status + internal network team
 - Restore + validate + postmortem
-

59. Explain autoscaling in real systems.

Answer: Scale based on CPU/RPS/queue depth, but ensure DB/cache layers can handle scaling too.

60. How do you ensure compliance in cloud infra?

Answer:

- Policy-as-code
 - Audit logs
 - Encryption
 - IAM least privilege
 - Automated compliance scans
 - Regular security reviews
-

61. Blue-green vs Canary—when to use which?

Answer:

- **Blue-green:** best for fast switch/rollback
- **Canary:** best for gradual safe rollout with real users

62. How do you ensure traceability in CI/CD pipelines?

Answer:

- Git commit tagging
 - Build IDs linked to deployments
 - Artifact versioning
 - Audit trails (who deployed what, when)
-

63. Network architecture change without disrupting production—approach?

Answer:

- Deploy changes in parallel
 - Use staged rollout
 - Validate with test traffic
 - Use feature flags and fallback routes
-

64. What is capacity planning and how do you do it?

Answer:

Use traffic trends, load testing, resource utilization history, and forecast growth to prevent future bottlenecks.

65. Intermittent connectivity between services—how do you debug?

Answer:

- Check DNS resolution
 - Network policies / SG rules
 - Packet loss metrics
 - Service discovery issues
 - Timeouts and retries
 - Trace request paths
-

66. Critical security patch needed everywhere—your execution plan?

Answer:

- Prioritize by severity
 - Patch staging first
 - Automate rollout using CI/CD
 - Use rolling updates
 - Validate monitoring
 - Document and confirm completion
-

67. How do you test a DR plan?

Answer:

By running scheduled DR drills: restore backups, failover traffic, validate application, verify RTO/RPO.

68. Cloud costs too high—what's your step-by-step approach?

Answer:

- Identify top cost services
 - Find idle resources
 - Right-size compute
 - Optimize storage
 - Use commitments (reserved instances)
 - Improve architecture (caching/CDN)
-

69. Multi-region redundancy strategy?

Answer:

- Active-passive or active-active
 - Global load balancing (Route53/Traffic manager)
 - Multi-region DB replication
 - Automated failover
-

70. Dependency/version management in CI/CD pipelines—how?

Answer:

- Lock versions
 - Use artifact repositories
 - Use semantic versioning
 - Automate dependency checks and upgrades
 - Maintain rollback compatibility
-

71. Integrating third-party APIs reliably—how?

Answer:

- Retries with exponential backoff
 - Circuit breaker pattern
 - Timeouts
 - Monitoring of external latency
 - Fallback responses
 - Rate limiting
-



**SRE QUICK REVISION CHEAT SHEET
(Interview Gold)**

Core Metrics

- Availability
- Latency (p95/p99)
- Error Rate
- Saturation (CPU/memory/disk/threads)
- Traffic (RPS)

Key Concepts

- SLI / SLO / SLA
- Error Budget
- Toil
- MTTR / MTBF
- Postmortem
- Observability (Metrics + Logs + Traces)

Golden Signals (Google SRE Standard)

✓ Latency • Traffic • Errors • Saturation

KUBERNETES INTERVIEW QUESTIONS + ANSWERS (28 Questions)

(Beginner → Intermediate → Advanced → Production Scenarios)

Kubernetes Interview Guide (28 Q&A)

Section A — Beginner Level (1–10)

1. What is Kubernetes and why is it used?

Answer: Kubernetes is an open-source container orchestration platform used to deploy, scale, and manage containerized applications automatically across a cluster of machines.

2. What is a Kubernetes Cluster?

Answer: A Kubernetes cluster is a group of machines (nodes) working together to run containerized applications. It consists of a **control plane** and **worker nodes**.

3. What is a Node in Kubernetes?

Answer: A node is a machine (VM or physical server) that runs application workloads. Nodes host **Pods**, and are managed by the control plane.

4. What is a Pod in Kubernetes?

Answer: A Pod is the smallest deployable unit in Kubernetes. It contains one or more containers sharing the same **network namespace** and **storage volumes**.

5. What is a Deployment in Kubernetes?

Answer: A Deployment manages a set of identical Pods and ensures the desired number of replicas are running. It supports **rolling updates** and **rollbacks**.

6. What is a ReplicaSet?

Answer: A ReplicaSet ensures a specific number of Pod replicas are running. Deployments use ReplicaSets internally.

7. What is a Service in Kubernetes?

Answer: A Service provides a stable network endpoint (DNS + IP) to expose Pods and load balance traffic across them.

8. What are the different types of Kubernetes Services?

Answer:

- **ClusterIP:** internal access only (default)
 - **NodePort:** exposes service on node's port
 - **LoadBalancer:** uses cloud provider LB
 - **ExternalName:** maps service to external DNS name
-

9. What is a Namespace in Kubernetes?

Answer: A Namespace is a logical separation within a cluster used to organize resources and manage access boundaries (dev, stage, prod).

10. What are labels and selectors in Kubernetes?

Answer:

- **Labels:** key-value metadata attached to resources
 - **Selectors:** used by Services/Deployments to match Pods by labels
-
-

Section B — Intermediate Level (11–20)

11. What is a ConfigMap?

Answer: A ConfigMap stores non-sensitive configuration data such as environment variables or config files, used by Pods at runtime.

12. What is a Secret in Kubernetes?

Answer: A Secret stores sensitive data like passwords, tokens, and keys. It can be mounted into Pods or exposed as environment variables.

13. What is the difference between ConfigMap and Secret?

Answer:

- ConfigMap = non-sensitive data
 - Secret = sensitive data (base64-encoded and managed more securely)
-

14. What is Ingress in Kubernetes?

Answer: Ingress is an API resource used to manage external HTTP/HTTPS access to services, typically with routing rules and TLS termination.

15. What is an Ingress Controller?

Answer: An Ingress Controller is the actual component (like NGINX Ingress Controller) that implements the Ingress rules and routes traffic.

16. What is the difference between a StatefulSet and a Deployment?

Answer:

- Deployment: for stateless workloads

- **StatefulSet:** for stateful apps (stable pod identity, stable storage)
-

17. What is a Persistent Volume (PV) and Persistent Volume Claim (PVC)?

Answer:

- **PV:** storage resource in the cluster
 - **PVC:** a request for storage by a Pod
PVC binds to PV based on size and access mode.
-

18. What are liveness and readiness probes?

Answer:

- **Liveness probe:** checks if container is alive (restart if fails)
 - **Readiness probe:** checks if container is ready to serve traffic
-

19. What is a DaemonSet in Kubernetes?

Answer: A DaemonSet ensures one Pod runs on each node (or selected nodes), commonly used for logging/monitoring agents.

20. What is a Job and a CronJob?

Answer:

- **Job:** runs a task until completion
 - **CronJob:** runs Jobs on a schedule (like Linux cron)
-
-

Section C — Advanced Level (21–24)

21. How does Kubernetes perform rolling updates?

Answer: Kubernetes gradually replaces old Pods with new Pods while maintaining availability. It uses `maxSurge` and `maxUnavailable` settings.

22. How do you rollback a failed deployment?

Answer:

You can rollback using:

```
kubectl rollout undo deployment <deployment-name>
```

It restores the previous ReplicaSet revision.

23. What is Horizontal Pod Autoscaler (HPA)?

Answer: HPA automatically scales the number of Pod replicas based on CPU/memory usage or custom metrics.

24. What is Cluster Autoscaler?

Answer: Cluster Autoscaler automatically scales worker nodes up/down when Pods can't be scheduled due to insufficient resources.

Section D — Production Scenario Questions (25–28)

25. A Pod is stuck in CrashLoopBackOff. How do you troubleshoot it?

Answer:

Steps:

1. Check logs: `kubectl logs <pod>`

2. Describe pod: `kubectl describe pod <pod>`
 3. Verify env variables, configmaps, secrets
 4. Validate image, command, and args
 5. Check probes (readiness/liveness causing restarts)
-

26. A Pod is Pending for a long time. What could be the reasons?

Answer:

Common causes:

- Insufficient CPU/memory in nodes
 - Node selectors/taints not matching
 - PVC not bound
 - Cluster Autoscaler not adding nodes
- Check using: `kubectl describe pod <pod>`
-

27. Service is not reachable from outside. What checks do you perform?

Answer:

- Service type correct? (NodePort / LoadBalancer / Ingress)
 - Ingress controller running?
 - DNS pointing correctly?
 - Security Groups / firewall open?
 - Pod labels match Service selectors?
-

28. High latency in production Kubernetes app. What's your approach?

Answer:

- Check resource saturation (CPU/memory throttling)
 - Check HPA scaling behavior
 - Analyze metrics + traces
 - Check DB latency and connection pool
 - Inspect networking (DNS, ingress, service mesh)
 - Add caching or optimize queries
-

Kubernetes Quick Revision (One-Liners)

- **Pod:** smallest unit
- **Deployment:** manages stateless pods
- **StatefulSet:** stable identity + storage
- **Service:** stable endpoint + load balancing
- **Ingress:** HTTP routing from outside
- **HPA:** scales pods
- **Cluster Autoscaler:** scales nodes
- **ConfigMap vs Secret:** config vs sensitive data
- **PV/PVC:** storage supply vs storage request

JENKINS INTERVIEW QUESTIONS + ANSWERS (28 Questions)

(Beginner → Intermediate → Advanced → Production Scenarios)

Jenkins Interview Guide (28 Q&A)

Section A — Beginner Level (1–10)

1. What is Jenkins and why is it used?

Answer: Jenkins is an open-source automation server used to implement **CI/CD pipelines**. It helps automate building, testing, and deploying software reliably.

2. What is the difference between CI and CD?

Answer:

- **CI (Continuous Integration):** frequent code merges + automated builds/tests
 - **CD (Continuous Delivery/Deployment):** automated delivery and release to environments (staging/production)
-

3. What is a Jenkins Job?

Answer: A Jenkins job is a task configured in Jenkins such as building code, running tests, or deploying an application.

4. What is Jenkins Pipeline?

Answer: A Jenkins Pipeline is a set of automated steps written as code to define how software is built, tested, and deployed.

5. What are the two types of Jenkins Pipelines?

Answer:

- **Declarative Pipeline:** structured and easier to read
 - **Scripted Pipeline:** flexible and uses Groovy scripting heavily
-

6. What is a Jenkinsfile?

Answer: A Jenkinsfile is a text file containing pipeline code, stored in the repository. It enables **Pipeline-as-Code** and version control.

7. What is a Jenkins Agent?

Answer: A Jenkins agent is a machine that runs builds and pipeline stages. Agents help distribute workloads and scale CI/CD execution.

8. What is a Jenkins Controller (Master)?

Answer: The Jenkins controller manages jobs, schedules builds, handles UI/API, stores configs, and coordinates with agents.

9. What is a build artifact in Jenkins?

Answer: A build artifact is an output file generated during pipeline execution, such as `.jar`, `.war`, Docker image tags, reports, etc.

10. What is the workspace in Jenkins?

Answer: Workspace is a directory on the Jenkins node where project files are checked out and build operations happen.

Section B — Intermediate Level (11–20)

11. How does Jenkins integrate with Git?

Answer: Jenkins connects to Git repositories via plugins and can automatically pull code using webhook triggers or polling.

12. What are Jenkins build triggers?

Answer: Build triggers define when jobs run, such as:

- SCM polling
 - Webhooks (GitHub/GitLab)
 - Scheduled (cron)
 - Upstream job completion
 - Manual triggers
-

13. What is the difference between Freestyle job and Pipeline?

Answer:

- **Freestyle:** UI-based, easier but limited and not scalable
 - **Pipeline:** code-driven, reusable, version-controlled, best for real CI/CD
-

14. What are Jenkins plugins and why are they important?

Answer: Plugins extend Jenkins functionality (Git, Docker, Kubernetes, Slack alerts, Maven, etc.). Jenkins is plugin-driven.

15. How do you securely store credentials in Jenkins?

Answer: Use **Jenkins Credentials Manager** to store secrets such as tokens, usernames/passwords, SSH keys, and API keys.

16. How do you use credentials inside a Jenkins pipeline?

Answer: Use `withCredentials()` to access secrets safely without exposing them in logs.

Example:

- Inject username/password
 - Inject secret text token
 - Inject SSH private key
-

17. What are Jenkins Shared Libraries?

Answer: Shared Libraries allow reusing common pipeline code across multiple repositories (for standard stages like build/test/deploy).

18. How do you handle environment variables in Jenkins?

Answer: You can set environment variables:

- globally in Jenkins
 - per-job configuration
 - in Jenkinsfile using `environment {}` block
-

19. What is a parallel stage in Jenkins Pipeline?

Answer: Parallel stages run multiple tasks at the same time to speed up pipelines, such as running tests for multiple modules in parallel.

20. How do you archive artifacts in Jenkins?

Answer: Jenkins stores artifacts using `archiveArtifacts` so they can be downloaded later from build history.

Example use cases:

- Build outputs
 - Reports
 - Logs
 - Deployment packages
-
-

Section C — Advanced Level (21–24)

21. How do you implement rollback in Jenkins pipelines?

Answer:

Rollback strategy depends on the deployment type:

- redeploy previous build artifact
- rollback Kubernetes deployment
- rollback Helm release
- use Blue/Green or Canary strategy

Rollback must be automated and tested like normal deployment.

22. How do you build Docker images using Jenkins?

Answer: Jenkins can build Docker images using:

- Docker CLI inside pipeline
- Docker plugin
- Kaniko (for Kubernetes environments)
- BuildKit

Typical flow: build → tag → scan → push to registry.

23. How do you integrate Jenkins with Kubernetes?

Answer: Jenkins can run dynamic agents as Kubernetes Pods using Jenkins Kubernetes Plugin. It improves scalability and reduces cost.

24. How do you secure Jenkins in production?

Answer: Key Jenkins security practices:

- enable authentication (SSO/LDAP/OAuth)
- role-based access control (RBAC)
- limit admin users
- restrict script approvals
- rotate credentials
- keep Jenkins/plugins updated
- disable anonymous access
- secure secrets using Vault (optional)

Section D — Production Scenario Questions (25–28)

25. A Jenkins job fails randomly (flaky builds). How do you fix it?

Answer:

Steps:

1. Identify pattern: agent-specific or global issue
 2. Check logs for resource/timeouts
 3. Verify dependency versions
 4. Add retry logic only for flaky stages
 5. Make build idempotent
 6. Stabilize test environment (mock dependencies if needed)
-

26. Your Jenkins pipeline is very slow. What do you optimize?

Answer:

- use parallel stages
 - reduce checkout depth (shallow clone)
 - cache dependencies (Maven, npm, pip)
 - use faster agents (auto-scaling)
 - reduce Docker build time via layer caching
 - split pipeline into stages with clear boundaries
-

27. Jenkins controller went down. How do you recover quickly?

Answer:

Best-practice recovery steps:

- restore from backup (JENKINS_HOME)
 - store Jenkins home on persistent storage
 - use Infrastructure-as-Code to recreate Jenkins
 - keep plugins version-locked
 - use external secrets storage
 - implement disaster recovery plan
-

28. Secrets got exposed in Jenkins logs. What do you do immediately?

Answer:

1. Rotate the exposed credentials immediately
 2. Remove logs or restrict access to them
 3. Enable masking of secrets
 4. Update pipeline to use `withCredentials()`
 5. Add security audit and access review
 6. Implement secret scanning in CI/CD
-



Jenkins Quick Revision (One-Liners)

- **Jenkinsfile:** pipeline as code

- **Controller:** manages scheduling and UI
- **Agent:** runs builds
- **Credentials Manager:** stores secrets safely
- **Shared Libraries:** reuse pipeline code
- **Artifacts:** stored build outputs
- **Triggers:** webhook, schedule, SCM polling
- **Security:** RBAC + updates + limited plugins

TERRAFORM INTERVIEW QUESTIONS + ANSWERS (28 Questions)

(Beginner → Intermediate → Advanced → Production Scenarios)

Terraform Interview Guide (28 Q&A)

Section A — Beginner Level (1–10)

1. What is Terraform?

Answer: Terraform is an **Infrastructure as Code (IaC)** tool used to provision and manage infrastructure across cloud providers (AWS, Azure, GCP) using declarative configuration files.

2. What is Infrastructure as Code (IaC)?

Answer: IaC is managing infrastructure using code instead of manual processes. It allows infrastructure to be:

 repeatable •  version-controlled •  automated •  auditable

3. Is Terraform declarative or imperative?

Answer: Terraform is **declarative**, meaning you describe the desired final state and Terraform figures out how to reach it.

4. What is a Terraform Provider?

Answer: A provider is a plugin that lets Terraform interact with APIs of platforms like AWS, Kubernetes, GitHub, etc.

Example providers:

- `aws`
 - `azurerm`
 - `google`
 - `kubernetes`
-

5. What is a Terraform Resource?

Answer: A resource is a component that Terraform creates and manages, such as:

- EC2 instance
 - S3 bucket
 - VPC
 - IAM role
 - Kubernetes deployment
-

6. What are Terraform Variables used for?

Answer: Variables make Terraform code reusable and flexible across environments like dev/stage/prod.

Example use cases:

- instance type
 - region
 - environment name
 - CIDR blocks
-

7. What are Terraform Outputs?

Answer: Outputs allow Terraform to print useful values after execution, such as:

- public IP address
 - DNS name
 - VPC ID
 - database endpoint
-

8. What is the purpose of `terraform init`?

Answer: It initializes Terraform by:

- downloading providers
 - setting up backend
 - preparing modules
 - creating `.terraform/`
-

9. What is the purpose of `terraform plan`?

Answer: It generates an execution plan that shows what Terraform will create, update, or destroy **before applying changes**.

10. What is the purpose of `terraform apply`?

Answer: It executes the plan and applies the changes to create or update infrastructure.

Section B — Intermediate Level (11–20)

11. What is a Terraform State file?

Answer: Terraform state (`terraform.tfstate`) stores the mapping between resources in your cloud and your configuration code.

Terraform uses state to track:

what exists • what changed • what to update next

12. Where should Terraform state be stored in real projects?

Answer: In production, use `remote state` such as:

- S3 + DynamoDB lock (AWS)
- Terraform Cloud
- Azure Storage
- GCS bucket

Remote state ensures collaboration and safety.

13. What is state locking and why is it important?

Answer: State locking prevents multiple users/pipelines from modifying infrastructure at the same time, avoiding corruption.

Example: DynamoDB table used for locking in AWS.

14. What is a Terraform Module?

Answer: A module is a reusable set of Terraform files grouped together, such as:

- VPC module
- EKS module
- S3 module
- ECS module

Modules improve maintainability and reuse.

15. What is the difference between `count` and `for_each`?

Answer:

- `count`: best for identical resources using numeric indexes
- `for_each`: best for unique resources using keys (maps/sets)

 `for_each` is generally safer for long-term changes.

16. What are Terraform Data Sources?

Answer: Data sources allow Terraform to **read existing infrastructure** without creating it.

Example: Fetch existing VPC ID, subnet IDs, AMI ID, etc.

17. What is the difference between `terraform refresh` and `plan`?

Answer:

- `refresh`: syncs state with real infra
- `plan`: compares config vs state and shows changes

(Note: In newer workflows, refresh behavior is handled as part of plan/apply.)

18. What is drift in Terraform?

Answer: Drift happens when real infrastructure changes outside Terraform (manual changes), causing mismatch between Terraform state and reality.

19. How do you detect and fix drift?

Answer:

- ✓ Detect drift: `terraform plan`
 - ✓ Fix drift: `terraform apply` (reconcile)
 - ✓ Prevent drift: enforce IaC-only changes using IAM policies and governance
-

20. What is the Terraform Backend?

Answer: Backend defines where state is stored and how state operations are performed.

Common backend example:

- S3 backend for state
 - DynamoDB for locking
-
-

Section C — Advanced Level (21–24)

21. How do you handle secrets in Terraform securely?

Answer: Best practices:

- avoid putting secrets in `.tfvars` in Git
 - use secret managers (AWS Secrets Manager / Vault)
 - mark values as `sensitive = true`
 - restrict state file access
 - encrypt remote state storage
-

22. What are Terraform Workspaces and when do you use them?

Answer: Workspaces allow using the same codebase for multiple environments by keeping separate state files.

Examples:

- `dev`
- `stage`
- `prod`

 Good for simple setups

 For large organizations, separate state and repos are often better.

23. What is the lifecycle block in Terraform?

Answer: The lifecycle block controls resource behavior.

Common lifecycle settings:

- `prevent_destroy = true` (protect critical resources)

- `ignore_changes` (ignore specific changes like tags)
-

24. How do you manage infrastructure dependencies in Terraform?

Answer: Terraform automatically builds dependencies through references.

Example:

- subnet depends on VPC
- EC2 depends on subnet/security group

You can also force dependency using:

- `depends_on`
-
-

Section D — Production Scenario Questions (25–28)

25. Terraform apply failed halfway. What happens and how do you recover?

Answer: Terraform may leave partial infrastructure created.

Recovery steps:

1. Check error logs
 2. Fix the issue (permissions/quota/provider bug)
 3. Run `terraform plan`
 4. Run `terraform apply` again
 5. If resource is broken, delete carefully or use `terraform taint` / recreate
-

26. Two engineers ran Terraform at the same time. What can break?

Answer: It may cause:

- corrupted state file
- conflicting resource updates
- unexpected deletions/overwrites

Fix: implement remote state + locking (example: DynamoDB lock).

27. Your Terraform state file got exposed. What do you do?

Answer: Immediate actions:

1. Treat it as a security incident
 2. Rotate exposed credentials/secrets immediately
 3. Restrict backend access (IAM)
 4. Encrypt state storage (S3 SSE-KMS)
 5. Move sensitive secrets to secret manager
 6. Enable logging/auditing
-

28. You want a safe and controlled Terraform deployment in CI/CD. How?

Answer: Recommended CI/CD approach:

- ✓ `terraform fmt` (format check)
- ✓ `terraform validate`
- ✓ `terraform plan` and store the plan artifact
- ✓ manual approval for production
- ✓ `terraform apply` with approval
- ✓ use remote state + locking
- ✓ restrict privileges using IAM roles



Terraform Quick Revision (One-Liners)

- **Terraform:** IaC tool for provisioning infra
- **Provider:** connects Terraform to cloud API
- **State:** mapping of code ↔ infra
- **Remote state:** collaboration + security
- **Locking:** prevents state corruption
- **Modules:** reusable infra building blocks
- **Drift:** manual infra changes outside Terraform
- **Workspaces:** multiple environments, separate state



ANSIBLE INTERVIEW QUESTIONS + ANSWERS (28 Questions)

(Beginner → Intermediate → Advanced → Production Scenarios)

Ansible Interview Guide (28 Q&A)

Section A — Beginner Level (1–10)

1. What is Ansible?

Answer: Ansible is an **agentless configuration management and automation tool** used for:
✓ configuration management • ✓ application deployment • ✓ orchestration • ✓ provisioning

It uses **SSH (Linux)** and **WinRM (Windows)** to connect to machines.

2. Why is Ansible called agentless?

Answer: Because Ansible does **not require any agent/software** to be installed on the target servers.

Only the control node needs Ansible installed.

3. What is an Ansible Control Node?

Answer: The **Control Node** is the machine where Ansible is installed and from which we execute automation tasks.

4. What is an Ansible Managed Node?

Answer: Managed nodes are the target machines/servers that Ansible configures using playbooks/modules.

5. What is an Inventory in Ansible?

Answer: Inventory is a list of servers (hosts) Ansible will manage.

It can be:

- **Static** (INI/YAML file)
 - **Dynamic** (AWS / Azure / Kubernetes inventory sources)
-

6. What is a Playbook in Ansible?

Answer: A playbook is a YAML file containing a set of tasks that defines automation steps.

A playbook runs:

Play → Tasks → Modules → Results

7. What is an Ansible Module?

Answer: Modules are reusable units of work used in tasks, such as:

- `yum`, `apt` (packages)
 - `copy`, `template` (files)
 - `service`, `systemd` (services)
 - `user` (users)
 - `shell`, `command` (commands)
-

8. What is the difference between `command` and `shell` module?

Answer:

- `command`: does **NOT** support shell features like pipes, redirects, variables
- `shell`: supports shell features like `|`, `&&`, redirection `>`

 Best practice: use `command` when possible (safer).

9. What are Ansible Facts?

Answer: Facts are system details automatically collected from managed nodes, like:

- OS type
- IP address
- CPU/memory
- hostname

Facts are stored inside `ansible_facts`.

10. What is idempotency in Ansible?

Answer: Idempotency means running the same playbook multiple times results in the **same state**, without unnecessary changes.

 Example: installing nginx repeatedly won't reinstall it again and again.

Section B — Intermediate Level (11–20)

11. What are Roles in Ansible?

Answer: Roles are a way to organize Ansible code into structured reusable components such as:

- tasks
- handlers
- templates
- vars
- defaults
- files

Roles make automation clean and reusable in projects.

12. What are Handlers in Ansible?

Answer: Handlers are tasks triggered only when a change happens.

Example: restart nginx only if config changes.

 Handlers run at the end of a play unless flushed explicitly.

13. What is Ansible Vault?

Answer: Ansible Vault is used to **encrypt secrets** like:

- passwords
- API keys
- certificates
- sensitive variables

 Prevents secret leakage in Git repositories.

14. What is the difference between Variables, Defaults, and Vars?

Answer:

- **defaults/** → lowest priority, safe fallback
- **vars/** → higher priority, overrides defaults
- Variables can also come from: inventory, CLI, facts, etc.

 This determines which variable wins.

15. What are Tags in Ansible?

Answer: Tags allow running specific tasks only.

Example use cases:

- run only deployment tasks
 - skip database tasks
 - run only “install” section
-

16. What is a Jinja2 template in Ansible?

Answer: Jinja2 templates allow inserting variables dynamically in config files.

Example: generate an Nginx config file based on environment values.

17. How do you run playbooks on only one host in a group?

Answer: You can limit execution using:

- `--limit hostname`
- or limit to a subgroup pattern

Example: run only on one server during testing.

18. What is the difference between `vars_files` and inventory variables?

Answer:

- `vars_files`: variables stored in separate YAML files
- inventory variables: variables directly attached to hosts/groups

Use inventory variables for environment-specific configuration.

19. What is `become` in Ansible?

Answer: `become` is used for privilege escalation (like sudo).

Example: install packages as root.

 Very common in Linux automation.

20. How do you check syntax and validate a playbook before running?

Answer:

- syntax check: `--syntax-check`
- dry-run simulation: `--check`
- show changes: `--diff`

 Prevents breaking production accidentally.

Section C — Advanced Level (21–24)

21. What is a Dynamic Inventory and why is it useful?

Answer: Dynamic inventory automatically fetches host details from sources like:

- AWS EC2 instances
- Kubernetes clusters
- Azure VMs
- CMDB tools

 Useful when servers scale up/down frequently.

22. What is the difference between `include_tasks` and `import_tasks`?

Answer:

- `import_tasks`: static, loaded at playbook parsing time
- `include_tasks`: dynamic, loaded during runtime

 Use `include` when conditions decide tasks at runtime.

23. How do you handle failures in Ansible playbooks?

Answer: Common strategies:

- `ignore_errors: yes` (only if safe)
 - `failed_when:` for custom failure logic
 - `retries:` and `until:` for retry loops
 - `block / rescue / always` for error recovery flows
-

24. How do you run tasks in parallel in Ansible?

Answer: Ansible runs tasks in parallel by default (based on forks).

You can tune performance using:

- `forks` settings
- `serial` for controlled batch rollout

 Useful for safe deployments (like 10 servers at a time).

Section D — Production Scenario Questions (25–28)

25. Scenario: A playbook works in staging but fails in production. What do you check?

Answer: Check common mismatches:

-  different OS version / packages
-  different permissions
-  firewall / DNS issues
-  missing environment variables

- secrets differences (Vault)
- inventory group variables incorrect

Fix by adding validation tasks, environment checks, and strict inventories.

26. Scenario: You ran a playbook and it broke a service. How do you rollback safely?

Answer: Rollback strategies:

- maintain backup configs before applying changes
- use versioned deployments (blue/green)
- use Ansible to restore previous config + restart service
- deploy in batches using `serial`

Always test rollback automation before production.

27. Scenario: You need to deploy updates with zero downtime using Ansible. How?

Answer: Options:

- rolling deployment (batch servers)
 - load balancer drain + update + reattach
 - health checks after each batch
 - `serial` and handlers to control restart timing
 - canary release to small percentage first
-

28. Scenario: How do you secure Ansible in real organizations?

Answer: Best practices:

- use Ansible Vault for secrets
- restrict SSH keys and sudo access

- use least privilege accounts
 - store playbooks in Git with review process
 - use centralized logging/audit trails
 - separate inventories per environment (dev/stage/prod)
-



Ansible Quick Revision (One-Liners)

- **Ansible:** automation & configuration management tool
- **Agentless:** no software needed on managed nodes
- **Inventory:** list of servers
- **Playbook:** YAML automation workflow
- **Module:** unit of work
- **Idempotent:** safe repeated runs
- **Handlers:** run only when notified
- **Vault:** encrypt secrets
- **Roles:** reusable structure
- **Tags:** run selected tasks only



KUBERNETES INTERVIEW QUESTIONS + ANSWERS (28 Questions)

(Beginner → Intermediate → Advanced → Production Scenarios)

Kubernetes Interview Guide (28 Q&A)

Section A — Beginner Level (1–10)

1. What is Kubernetes and why is it used?

Answer: Kubernetes (K8s) is an **open-source container orchestration platform** used to:

- deploy containerized applications
 - scale applications automatically
 - manage self-healing (restart, reschedule)
 - perform rolling updates and rollbacks
 - handle service discovery & load balancing
-

2. What problem does Kubernetes solve?

Answer: Kubernetes solves problems such as:

- managing large numbers of containers
- maintaining application availability
- scaling services reliably
- automating deployments and infrastructure coordination

Without Kubernetes, managing containers across multiple servers becomes manual and error-prone.

3. What is a Kubernetes Cluster?

Answer: A Kubernetes cluster is a group of machines that run containerized workloads.

It includes:

- Control Plane (Master Components)**
 - Worker Nodes (where apps run)**
-

4. What is a Node in Kubernetes?

Answer: A node is a machine (VM or physical server) that runs application workloads.
Each node contains:

- **kubelet** (node agent)
 - **container runtime** (Docker/containerd)
 - **kube-proxy** (network rules)
-

5. What is a Pod?

Answer: A pod is the **smallest deployable unit in Kubernetes**.
It can contain:

- **one container** (most common)
- or **multiple containers** that share:
 - network (same IP)
 - storage volumes

 Pods are ephemeral (can be destroyed and recreated).

6. What is the difference between a Pod and a Container?

Answer:

- A **container** is a runtime package of an application
- A **pod** is a Kubernetes wrapper for containers

 A pod can run **one or more containers**, while a container is a single unit.

7. What is a Deployment in Kubernetes?

Answer: A Deployment is used to manage application pods declaratively. It provides:

- desired number of replicas
 - rolling updates
 - rollbacks
 - self-healing replacement pods
-

8. What is a ReplicaSet?

Answer: A ReplicaSet ensures the required number of pod replicas are running at all times.

- Deployment manages ReplicaSets automatically.
-

9. What is a Kubernetes Service?

Answer: A Service provides a stable way to access pods, because pod IPs change frequently.

Types of Services:

- **ClusterIP** (internal only)
 - **NodePort** (expose on worker node port)
 - **LoadBalancer** (cloud load balancer)
 - **ExternalName** (DNS mapping)
-

10. What is a Namespace in Kubernetes?

Answer: A namespace is a logical separation inside a Kubernetes cluster used to:

- isolate teams/projects
- manage resources separately
- apply RBAC and policies differently

Common namespaces: `default`, `kube-system`, `kube-public`.

Section B — Intermediate Level (11–20)

11. What are Labels and Selectors in Kubernetes?

Answer:

- **Labels:** key-value metadata attached to objects (`app=web`)
- **Selectors:** used to match objects based on labels

 Services use selectors to connect to pods.

12. What is a ConfigMap?

Answer: A ConfigMap stores **non-sensitive configuration data** like:

- environment variables
- application configs
- URLs, flags, settings

 Keeps configuration separate from container image.

13. What is a Secret in Kubernetes?

Answer: Secrets store **sensitive data**, such as:

- DB passwords
- API keys
- TLS certificates

 Best practice: do not hardcode secrets in YAML or Git.

14. What are Liveness and Readiness Probes?

Answer:

- **Liveness Probe:** checks if the container is alive → restarts if failed
- **Readiness Probe:** checks if the container is ready to serve traffic → removes from service if failed

They improve reliability and zero-downtime deployments.

15. What is a StatefulSet and when do you use it?

Answer: StatefulSet is used for **stateful applications** requiring:

- stable pod identity (fixed name)
- stable storage
- ordered startup/shutdown

Examples: MySQL, MongoDB, Kafka, Elasticsearch.

16. What is the difference between Deployment and StatefulSet?

Answer:

- Deployment:** stateless workloads, random pod names, easy scaling
 - StatefulSet:** stable identity + stable storage, ordered rolling operations
-

17. What is a Persistent Volume (PV) and Persistent Volume Claim (PVC)?

Answer:

- **PV:** actual storage resource in cluster
- **PVC:** request for storage by a pod

Pod uses PVC → PVC binds to PV.

18. What is Ingress in Kubernetes?

Answer: Ingress provides **HTTP/HTTPS routing** to services, allowing:

- domain-based routing
- path-based routing
- TLS termination (HTTPS)

Ingress needs an **Ingress Controller** (Nginx, ALB, Traefik).

19. What is a DaemonSet?

Answer: DaemonSet ensures **one pod runs on every node** (or selected nodes).

Common use cases:

- logging agents (Fluentd)
 - monitoring agents (Node Exporter)
 - security agents
-

20. What is a Job and CronJob?

Answer:

- **Job:** runs a task until completion (one-time batch job)
- **CronJob:** runs jobs on a schedule (like Linux cron)

Examples: backups, cleanup scripts, scheduled reports.

Section C — Advanced Level (21–24)

21. What is the Kubernetes Control Plane?

Answer: The control plane manages the cluster and scheduling. Components include:

- **API Server** (entry point)
- **Scheduler** (places pods on nodes)
- **Controller Manager** (maintains desired state)
- **etcd** (cluster database)

If control plane fails, workloads may run but cluster becomes unmanaged.

22. What is etcd and why is it important?

Answer: etcd is a distributed key-value store used to store:

- cluster configuration
- secrets (unless external)
- state of nodes, pods, services

etcd must be backed up regularly in production.

23. What is HPA (Horizontal Pod Autoscaler)?

Answer: HPA automatically scales pods based on metrics such as:

- CPU usage
- memory usage
- custom metrics (via Prometheus adapter)

Useful for handling load spikes dynamically.

24. What is a NetworkPolicy?

Answer: NetworkPolicy controls network traffic rules between pods, namespaces, and services.

Provides micro-segmentation and zero-trust networking inside clusters.

Example: allow only backend pods to talk to database pods.

Section D — Production Scenario Questions (25–28)

25. Scenario: Pods are stuck in CrashLoopBackOff. What do you do?

Answer: Steps to troubleshoot:

- check logs (`kubectl logs`)
- describe the pod (`kubectl describe pod`)
- verify env vars, secrets, ConfigMaps
- check image name/tag and pull permissions
- check probes (liveness/readiness failures)
- check resource limits (OOMKilled)

Most common causes: missing configs, wrong command, app crash, memory limit.

26. Scenario: A deployment is slow and keeps restarting pods. How do you fix it?

Answer:

- verify readiness probes are correct
- increase initialDelaySeconds if app needs time
- increase CPU/memory resources
- check external dependencies (DB, API)
- enable autoscaling if needed
- check node capacity and scheduling delays

Readiness probe misconfiguration is one of the most common reasons.

27. Scenario: You need zero downtime deployment for a production app. What's your approach?

Answer:

- use rolling updates with proper probes
- set maxUnavailable=0 and maxSurge=1 (if needed)
- ensure graceful shutdown using:
 - preStop hooks
 - terminationGracePeriodSeconds
 - use Ingress/Service stable routing
 - perform canary release if high risk

28. Scenario: How do you secure a Kubernetes cluster in production?

Answer: Strong production security practices:

- RBAC (least privilege)
 - namespace isolation
 - network policies
 - secrets encryption + external secret managers
 - disable anonymous API access
 - restrict cluster-admin usage
 - scan images (Trivy, etc.)
 - use Pod Security Standards (restricted/baseline)
 - audit logs enabled
-

Kubernetes Quick Revision (One-Liners)

- **Pod:** smallest unit, runs container(s)
- **Deployment:** manages stateless pods + rollouts
- **StatefulSet:** stable identity + storage
- **Service:** stable access to pods
- **Ingress:** routing via HTTP/HTTPS

- **PV/PVC:** storage management
- **ConfigMap:** non-sensitive config
- **Secret:** sensitive config
- **HPA:** autoscale pods
- **NetworkPolicy:** restrict pod traffic

PYTHON + SHELL INTERVIEW QUESTIONS + ANSWERS (28 Questions)

(DevOps Focused | Beginner → Intermediate → Advanced → Production Scenarios)

Python + Shell Interview Guide (28 Q&A)

Section A — Beginner Level (1–10)

1. What is the main difference between Shell scripting and Python in DevOps?

Answer:

- **Shell scripting** is best for **quick OS-level automation** (files, services, commands, pipelines).
- **Python** is best for **complex automation**, API calls, data parsing, and scalable scripts.

 Rule of thumb:

Shell = glue automation | Python = structured automation

2. When should you choose Shell over Python?

Answer: Use Shell when the task involves:

- system administration commands (grep, awk, sed)
 - managing services (systemctl, ps, kill)
 - automating Linux tasks quickly
 - chaining commands via pipes
-

3. When should you choose Python over Shell?

Answer: Use Python when you need:

- API integration (AWS, GitHub, monitoring APIs)
 - JSON/YAML parsing
 - handling complex logic and error handling
 - better maintainability and reusable modules
-

4. What is the purpose of `#!/bin/bash` in a shell script?

Answer: It's called a **shebang**. It tells the system which interpreter should execute the script.

Example:

`#!/bin/bash` → run using Bash shell

5. How do you make a shell script executable?

Answer:

Use:

`chmod +x script.sh`

Then run:

`./script.sh`

6. What are environment variables and why are they important?

Answer: Environment variables store system values used by processes, such as:

- `PATH`

- HOME
- USER
- AWS_REGION

Used in DevOps for configuration without hardcoding.

7. How do you print a variable in Shell and Python?

Answer:

Shell:

```
echo $VAR
```

Python:

```
print(var)
```

8. How do you handle user input in Shell and Python?

Answer:

Shell:

```
read name
echo "Hello $name"
```

Python:

```
name = input("Enter name: ")
print("Hello", name)
```

9. How do you check exit status in Shell?

Answer:

In Shell, `$?` gives exit code of last command:

- `0` → success
- `non-zero` → failure

Example:

```
ls /tmp  
echo $?
```

10. What are the most common Shell commands used in DevOps daily work?

Answer: Some essentials are:

- `ls, cd, pwd`
 - `cp, mv, rm, mkdir`
 - `cat, tail, head, grep`
 - `ps, top, kill`
 - `systemctl, journalctl`
 - `curl, wget, ssh, scp`
-



Section B — Intermediate Level (11–20)

11. Explain pipes (`|`) in Shell scripting with an example.

Answer: Pipes send output of one command as input to another.

Example:

```
cat app.log | grep ERROR | wc -l
```

counts total ERROR lines in a log file.

12. What is the difference between `>` and `>>` in Shell?

Answer:

- `>` overwrites file
- `>>` appends to file

Example:

```
echo "Hello" > file.txt
echo "World" >> file.txt
```

13. What are `grep`, `awk`, and `sed` used for?

Answer:

- `grep` → search text
- `awk` → text processing + column extraction
- `sed` → stream editor (replace/edit text)

Example:

```
awk '{print $1}' file.txt
```

prints first column.

14. What is a function in Shell and Python?

Answer:

Shell function example:

```
hello() {  
    echo "Hello DevOps"  
}  
hello
```

Python function example:

```
def hello():  
    print("Hello DevOps")  
hello()
```

15. How do you run a command in background and keep it running after logout?

Answer:

run in background:

```
command &
```

keep running after logout:

```
nohup command &
```

16. How do you find which process is using a specific port?

Answer:

Use:

```
sudo lsof -i :8080
```

or

```
sudo netstat -tulnp | grep 8080
```

17. How do you parse JSON in Python for DevOps tasks?

Answer: Use Python's `json` module:

```
import json

data = json.loads(' {"env": "prod", "region": "us-east-1"}')
print(data["region"])
```

Very useful in cloud automation.

18. What is the difference between `subprocess` and `os.system()` in Python?

Answer:

- `os.system()` → simple but limited output handling
- `subprocess` → better control, captures output, safer

Example:

```
import subprocess
result = subprocess.run(["ls", "-l"], capture_output=True, text=True)
print(result.stdout)
```

19. How do you handle exceptions in Python?

Answer: Using `try-except`:

```
try:
    x = 10 / 0
except Exception as e:
    print("Error:", e)
```

 Important in production scripts to avoid crashes.

20. How do you schedule scripts in Linux?

Answer: Use cron jobs.

Example:

```
crontab -e
```

Run script every day at 2 AM:

```
0 2 * * * /path/backup.sh
```

Section C — Advanced Level (21–24)

21. What is idempotency and why does it matter in DevOps automation?

Answer: Idempotency means running a script multiple times produces the same final result.

 Example:

A script that creates a directory should not fail if the directory already exists.

Shell best practice:

```
mkdir -p /opt/app
```

22. How do you write a safe Bash script for production usage?

Answer: Recommended safety settings:

```
set -euo pipefail
```

Meaning:

- `-e` → stop on error
- `-u` → fail on undefined variables
- `pipefail` → fail if any command in pipeline fails

 Prevents silent failures.

23. Explain logging best practices for automation scripts.

Answer: Best practices:

-  timestamp logs
-  log to file and console
-  include error codes
-  log structured messages (INFO/WARN/ERROR)

Shell example:

```
echo "$(date) [INFO] Starting deployment..."
```

Python example using logging module:

```
import logging
logging.basicConfig(level=logging.INFO)
logging.info("Starting deployment")
```

24. How do you securely manage secrets in scripts?

Answer: Avoid hardcoding secrets inside scripts. Instead:

-  use environment variables
-  use secret managers (AWS Secrets Manager, Vault)
-  restrict permissions
-  rotate credentials

Bad example:

```
PASSWORD="admin123"
```

Good example:

```
PASSWORD=\"$DB_PASSWORD\"
```

Section D — Production Scenario Questions (25–28)

25. Scenario: A production server is running out of disk space. What do you do?

Answer:

 check disk usage:

```
df -h
```

 find heavy directories:

```
du -sh /* | sort -h
```

 check logs:

```
journalctl --disk-usage
```

 cleanup safely:

- remove old logs
- rotate logs
- remove unused docker images:

```
docker system prune
```

 prevent recurrence: monitoring + alerts.

26. Scenario: Your script works manually but fails in Jenkins/CI pipeline. Why?

Answer: Common causes:

- missing environment variables in CI
- permission issues (non-root agents)
- different PATH in CI
- missing files (workspace not correct)
- script requires interactive input

Fix tips:

- always define full paths
 - print env values inside pipeline for debugging
 - use non-interactive flags (like `-y`)
-

27. Scenario: A deployment script accidentally brings down production. How do you prevent this in the future?

Answer:

- implement safeguards:

- confirm environment (prod/staging)
- add approvals for prod deploy
- use feature flags
- use blue/green or canary deployments
- use rollback support
- add automated tests before deployment

- Most failures come from missing validation checks.
-

28. Scenario: You need to automate log parsing and alerting. How would you do it?

Answer: Approach:

- ✓ Shell quick parsing:

```
grep "ERROR" app.log | tail -20
```

- ✓ Python automation for better logic:

- parse logs
- count frequency
- send alert via email/Slack webhook
- run as cronjob

Python idea (high level):

- read file
- filter errors
- if threshold crossed → notify

- ✓ Python is ideal for scalable alert automation.
-

✓ Python + Shell Quick Revision Sheet

Shell Must-Know

- `grep, awk, sed`
- `ps, top, kill`

- `df -h`, `du -sh`
- pipes `|`, redirection `>` `>>`
- cron jobs
- `set -euo pipefail`

Python Must-Know

- `try/except`
- `subprocess`
- `json` parsing
- logging module
- scripts for API automation